



DD2437 – Artificial Neural Networks and Deep Architectures (annda)

Lecture 9: **Deep learning fundamentals**

General philosophy and a review of deep architectures

Pawel Herman

Computational Science and Technology (CST)

KTH Royal Institute of Technology

February 2018

AI ambition behind Deep Learning

The grand plan is to *“allow computers to model our world well enough to exhibit what we call intelligence”*.

(Bengio, 2006)

- The need for capturing high-level of abstraction
- Hope in learning algorithms that can help to exploit large quantities of available information (big data in the future) and generalise it to new contexts
- The assumption about the need for highly nonlinear (varying) mathematical functions (accounting for variations in the multivariate, often high-dimensional, domain of interest) to model complex behaviours

AI ambition behind Deep Learning

The grand plan is to *“allow computers to model our world well enough to exhibit what we call intelligence”*.

So, we need

- knowledge
- learning
 - complex functions,
 - from unlabeled data
 - with little human input
- generalisation
- understanding/identifying the underlying explanatory factors

AI ambition behind Deep Learning

The grand plan is to *“allow computers to model our world well enough to exhibit what we call intelligence”*.

So, we need

- knowledge
- learning
 - complex functions,
 - from unlabeled data
 - with little human input
- generalisation
- understanding/identifying the underlying explanatory factors



Is there a hope in
DEEP LEARNING in
DEEP ARCHITECTURES ?


What is depth in ML?

- **Depth of architecture**

- the number of levels of composition of nonlinear operations in the function learnt
- the length of the longest path from input to output in the graph

- **Deep learning**


- using multiple layers of inf. processing stages in hierarchical architectures for pattern recognition and representation learning
- focus on (incremental) learning of feature hierarchies



Motivation for deep structures

Why go deep? Do we need deep structures?

- Expressive power and compactness of models (*expressibility* and *efficiency*)
 - enhances generalisation, especially with limited training examples
 - less degrees of freedom when handling complexity and nonlinearity – exponential gain




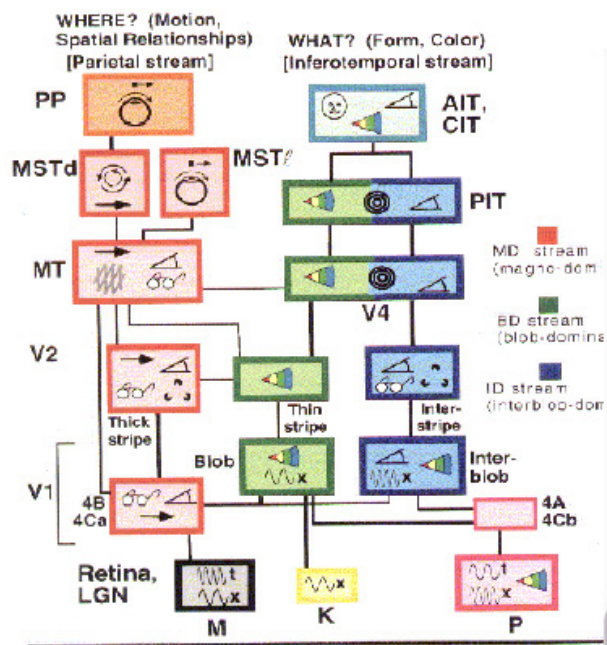
Shallow structure may need exponential size of hidden layer(s)

The universal approximation theorem and approximation costs.

Motivation for deep structures

Why go deep? Do we need deep structures?

- Inspirations from hierarchical brain organisation



LeCun & Ranzato, 2013

Motivation for deep structures

Why go deep? Do we need deep structures?



- Expressive power and compactness of models
 - enhances generalisation, especially with limited training examples
 - less degrees of freedom when handling complexity and nonlinearity
- Inspirations from hierarchical brain organisation
- Cognitive inspiration – multiple levels of abstraction

Motivation for deep structures

Why go deep? Do we need deep structures?

- Finally,


multiple levels of representations facilitate transfer and multi-task learning (hierarchy of representations, non-local generalisation)



Lee, 2011

Trouble with classical multi-layer ANNs


- Hard to train
 - the problem of vanishing gradients (diffusion of gradients) in backpropagation algorithm
 - non-convex optimisation
 - local minima
 - susceptibility to overfitting



Learning high-level features – data representations

Traditional pattern recognition

- Human-designed representations (hand-engineered features)




- Focus on optimisation to make best predictions
- Importance of data labels in supervised learning

(\mathbf{x}, \mathbf{y})

Learning high-level features – data representations

Traditional pattern recognition

- Human-designed representations (hand-engineered features)




- Focus on optimisation to make best predictions
- Importance of data labels in supervised learning

(x, y)

Deep learning approach


- Representation learning where good features are automat. learnt
- Potential to learn multiple levels of representation in DL algorithms



Learning high-level features – data representations

Traditional pattern recognition

- Human-designed representations (hand-engineered features)



- Focus on optimisation to make best predictions
- Importance of data labels in supervised learning

(\mathbf{x}, \mathbf{y})


Deep learning approach

- Representation learning where good features are automat. learnt
- Potential to learn multiple levels of representation in DL algorithms
- For example,
 - character → word → word group, phrase → clause → sentence → story
 - pixel → edge → motif → object
 - sample → spectral feature → sound → phoneme → word

Learning high-level features – data representations

Traditional pattern recognition

- Human-designed representations (hand-engineered features)



- Focus on optimisation to make best predictions
- Importance of data labels in supervised learning

(\mathbf{x}, \mathbf{y})

Deep learning approach

- Representation learning where good features are automat. learnt
- Potential to learn multiple levels of representation in DL algorithms
- Good predictions are v. important but so is data representation
- Both unsupervised and supervised mode is heavily exploited – unlabeled data are also useful

Short historical note on deep architectures in ML (1)

- Perceptron - the first learning machine (Rosenblatt, ~1960)
- Deep learning in artificial neural networks
 - revival of interest with backpropagation in 1980s
 - "better" backprop with advanced gradient descent
 - generalisation –complexity issues and bias/variance dilemma

Short historical note on deep architectures in ML (1)

- Perceptron - the first learning machine (Rosenblatt, ~1960)
 - Deep learning in artificial neural networks
 - revival of interest with backpropagation in 1980s
 - "better" backprop with advanced gradient descent
 - generalisation –complexity issues and bias/variance dilemma
- BUT still.....**
- lack of ability to learn from the unlabeled data (most data is unlabeled)
 - slow learning, problems with convergence, sensitivity to local minima

Short historical note on deep architectures in ML (2)

- Perceptron - the first learning machine (Rosenblatt, ~1960)
- Deep learning in artificial neural networks
 - revival of interest with backpropagation in 1980s
 - "better" backprop with advanced gradient descent
 - generalisation –complexity issues and bias/variance dilemma
- Shallow architectures with Support Vector Machines (SVMs)
 - effective in addressing simple and well-constrained problems
 - kernels arbitrarily define features (not hand-crafted but still "fixed")
 - limited modelling and representational power

Short historical note on deep architectures in ML (2)

- Perceptron - the first learning machine (Rosenblatt, ~1960)
- Deep learning in artificial neural networks
 - revival of interest with backpropagation in 1980s
 - "better" backprop with advanced gradient descent
 - generalisation –complexity issues and bias/variance dilemma
- Shallow architectures with Support Vector Machines (SVMs)
 - effective in addressing simple and well-constrained problems
 - kernels arbitrarily define features (not hand-crafted but still "fixed")
 - limited modelling and representational power

Prior knowledge is
arbitrary, not learnt

Short historical note on deep architectures in ML (3)

- Exploration of potential benefits of unsupervised or semi-supervised techniques in the context of supervised learning

Short historical note on deep architectures in ML (3)

- Exploration of potential benefits of unsupervised or semi-supervised techniques in the context of supervised learning
- Identification of *Fundamental Deep Learning Problem* in 1991
 - vanishing or exploding gradients – **unstable learning**
 - progressive ideas with deep hierarchies of recurrent networks, auto-encoders and first deep belief networks

Short historical note on deep architectures in ML (3)

- Exploration of potential benefits of unsupervised or semi-supervised techniques in the context of supervised learning
- Identification of *Fundamental Deep Learning Problem* in 1991
 - vanishing or exploding gradients – **unstable learning**
 - progressive ideas with deep hierarchies of recurrent networks, auto-encoders and first deep belief networks
- Major breakthrough in 2006
 - the idea to pre-train deep architectures with layer-wise unsupervised learning (groups led by G.E. Hinton, Y. Bengio and Y. LeCun)
 - more efficient parameter estimation methods

Short historical note on deep architectures in ML (3)

- Exploration of potential benefits of unsupervised or semi-supervised techniques in the context of supervised learning
- Identification of *Fundamental Deep Learning Problem* in 1991
 - vanishing or exploding gradients – **unstable learning**
 - progressive ideas with deep hierarchies of recurrent networks, auto-

[1] Hinton, G. et al. (2006) A fast learning algorithm for deep belief nets. *Neural Computation* 18:1527-1554,

[2] Bengio, Y. et al. (2006) Greedy Layer-Wise Training of Deep Networks, in J. Platt et al. (Eds), *Advances in Neural Information Processing Systems 19 (NIPS 2006)*, pp. 153-160.

[3] Ranzato, M. et al. & Yann LeCun, Y. (2006) Efficient Learning of Sparse Representations with an Energy-Based Model, in J. Platt et al. (Eds), *NIPS*.

Short historical note on deep architectures in ML (3)

- Exploration of potential benefits of unsupervised or semi-supervised techniques in the context of supervised learning
- Identification of *Fundamental Deep Learning Problem* in 1991
 - vanishing or exploding gradients – **unstable learning**
 - progressive ideas with deep hierarchies of recurrent networks, auto-

Shared principles in these papers:


- - Unsupervised learning of representations is used to (pre-)train each layer.
 - Unsupervised training of one layer at a time, on top of the previously trained ones. The representation learned at each level is the input for the next layer.
 - Use supervised training to fine-tune all the layers (in addition to one or more additional layers that are dedicated to producing predictions).

Successful applications as a driver

- Convolutional nets (CNNs) in *computer vision*
- Deep learning based *speech recognition* systems developed by Google and Microsoft
- Deep learning is becoming a hot topic in *natural language processing* (NLP)
- Advances in machine translation (RNNs, LSTM)
- Growing importance in reinforcement learning (deep RL)
- Scope of applications massively grows

General theme of the older deep learning protocol – deep belief networks, stacked autoencoders

- Greedy layer-wise **unsupervised pre-training** +
supervised tuning (the legacy of Hinton, Bengio and LeCun)

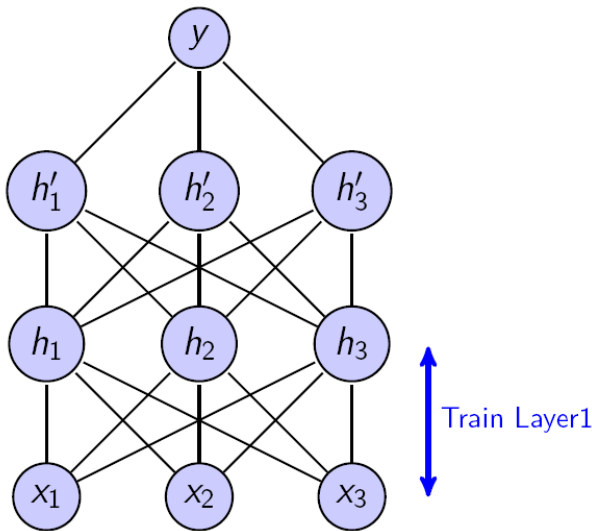


Single layer at a time


Hinton et al., 2006
Duh, 2013

General theme of the older deep learning protocol – deep belief networks, stacked autoencoders

- Greedy layer-wise **unsupervised pre-training** + supervised tuning (the legacy of Hinton, Bengio and LeCun)



Single layer at a time

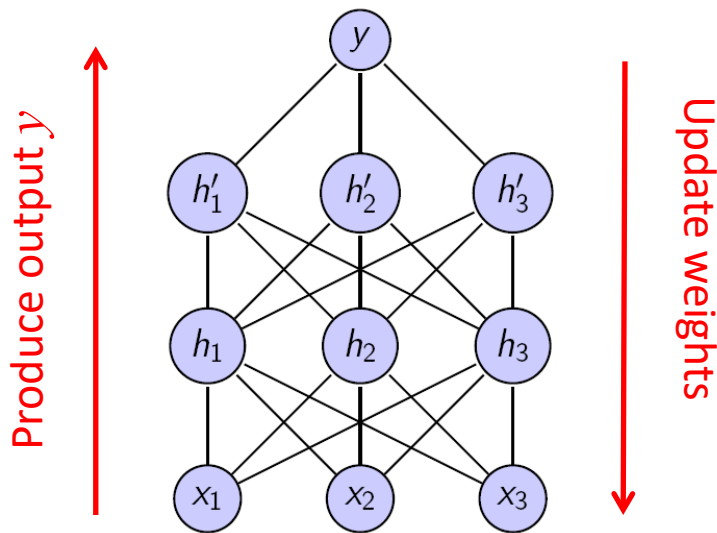


Train another layer while keeping the lower layer fixed

Hinton et al., 2006
Duh, 2013

General theme of the older deep learning protocol – deep belief networks, stacked autoencoders

- Greedy layer-wise unsupervised pre-training +
supervised tuning (the legacy of Hinton, Bengio and LeCun)



Gradient-based fine tuning

1. Add a classifier layer and retrain globally the entire structure.
2. Train only a supervised classifier on top and keep other layers fixed.

Hinton et al., 2006
Duh, 2013
LeCun & Ranzato, 2013

Hypothetical role of unsupervised pre-training

- Regularisation hypothesis (Erhan et al., 2010)
 - Pre-training minimises **variance**
 - It also helps to control **complexity** for architectures with large sizes of hidden layers
 - Acts like an implicit penalisation term – **regularisation**


Hypothetical role of unsupervised pre-training

- Regularisation hypothesis (Erhan et al., 2010)
 - Pre-training minimises variance
 - It also helps to control complexity for architectures with large sizes of hidden layers
 - Acts like an implicit penalisation term – regularisation
- Optimisation hypothesis (Bengio et al., 2007)



Hypothetical role of unsupervised pre-training


- Regularisation hypothesis (Erhan et al., 2010)
 - Pre-training minimises variance
 - It also helps to control complexity for architectures with large sizes of hidden layers
 - Acts like an implicit penalisation term
- Optimisation hypothesis (Bengio et al., 2007)
 - pre-training finds a better initial condition for further gradient-based optimisation
 - it facilitates training of the entire architecture (lower and higher layers benefit from tuning)



Most common network architecture and learning types


Restricted Boltzmann machine (RBM) layer

(contrastive divergence for pre-training)




Auto-encoder (AE) layer

(gradient descent based algorithms for pre-training)




Convolutional neural networks (CNNs)



Greedy layer-wise unsupervised pre-training, which is increasingly omitted once **ReLU** units are employed


Network initialised without any pre-training

Restricted Boltzmann machine (RBM)



In traditional RBM, x_i and h_j are binary variables

Simple energy-based model




$$p(x, h) \sim e^{-E_{\theta}(x, h)}$$

$$E_{\theta}(x, h) = -x'Wh - b'x - d'h$$

The idea is to optimise log-likelihood with the use of approximative Gibbs sampling – **Contrastive Divergence** algorithm

Auto-encoders



Decoder: $x' = \sigma(W'h + d)$


Encoder: $h = \sigma(Wx + b)$

Encourage h to give small reconstruction error:

- e.g. $Loss = \sum_m \|x^{(m)} - DECODER(ENCODER(x^{(m)}))\|^2$
- Reconstruction: $x' = \sigma(W'\sigma(Wx + b) + d)$



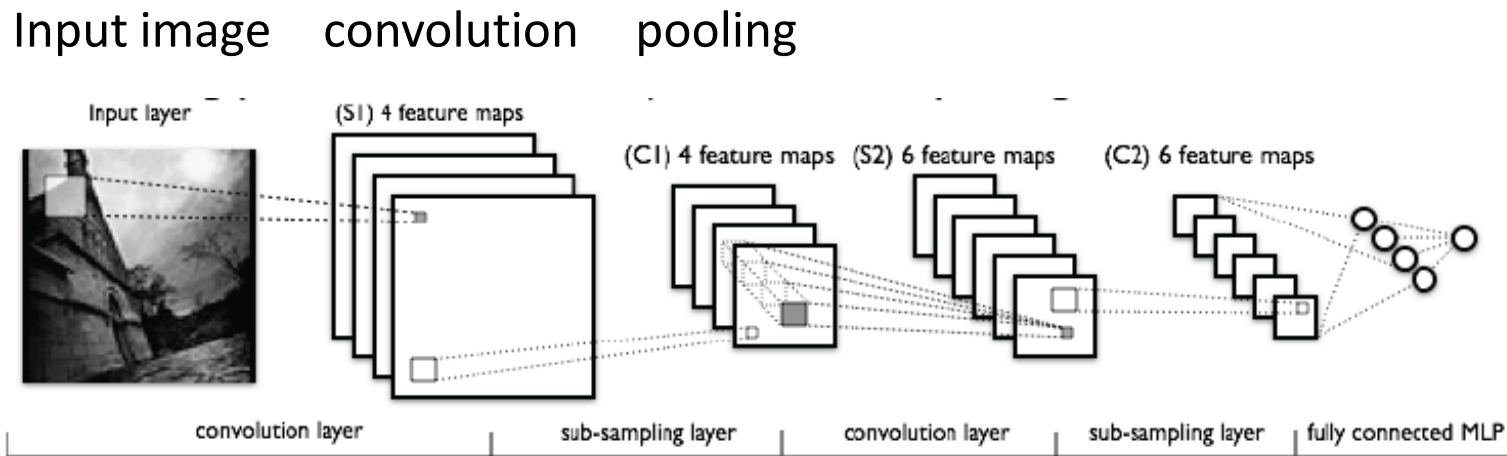
(REF)

Convolutional neural networks (CNNs)



LeCun et al., 1989

Convolutional neural networks (CNNs)



LeCun et al., 1989

Generative vs discriminative approach

1. Generative deep architectures

- describe statistical distributions of data and associated classes, $P(X,Y)$
- characterise higher-order correlational structure of data for pattern analysis (suitable for holistic training of complex systems)
- energy-based models including auto-encoders

Generative vs discriminative approach

1. Generative deep architectures

- describe statistical distributions of data and associated classes, $P(X,Y)$
- characterise higher-order correlational structure of data for pattern analysis (suitable for holistic training of complex systems)
- energy-based models including auto-encoders


2. Discriminative deep architectures

- provide discriminative power for pattern classification by characterising the posterior distribution $P(Y|X)$
- HMM, CNN, DBN-DNN

Generative vs discriminative approach

3. Hybrid deep architectures


- the goal is discrimination but is helped by the outcomes of generative modelling in deep architectures
- at the heart of early ideas for deep learning proposed by Hinton, Bengio and LeCun – unsupervised learning + supervised tuning
- deep belief networks (DBNs) are considered as a precursor component of hybrid deep architectures.



Deng, 2013


Why should we be bothered with deep learning?

- Learning representations (-> more in the next lecture)
 - learning features as part of DL algorithms
 - multiple levels of abstraction and complexity (hierarchy)



Why should we be bothered with deep learning?

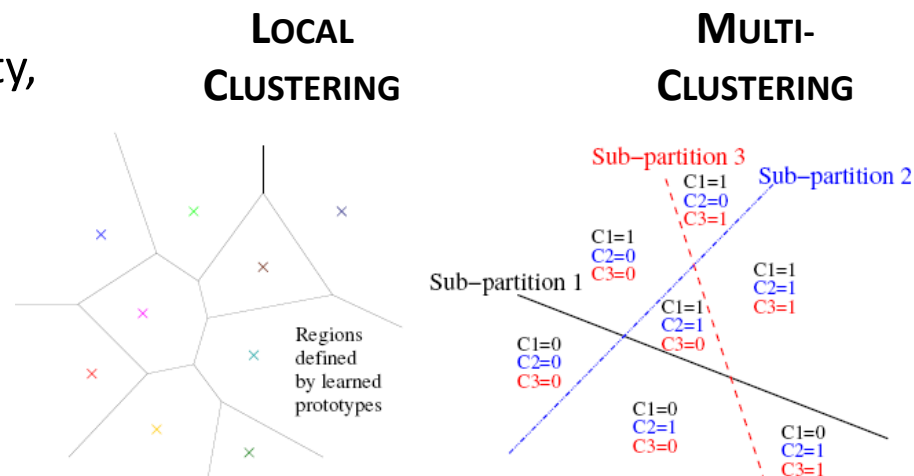
- Learning representations
 - learning features as part of DL algorithms
 - multiple levels of abstraction and complexity (hierarchy)
- Distributed feature representations
 - multi-task or transfer learning



Why should we be bothered with deep learning?

- Learning representations
 - learning features as part of DL algorithms
 - multiple levels of abstraction and complexity (hierarchy)
- Distributed feature representations
 - multi-task or transfer learning
 - mitigates the curse of dimensionality, allows for non-local generalisation (multi-clustering)
 - sparse coding

Bengio and Delalleau, 2013



Why should we be bothered with deep learning?

- Learning representations
 - learning features as part of DL algorithms
 - multiple levels of abstraction and complexity (hierarchy)
- Distributed feature representations
 - multi-task or transfer learning (multi-clustering)
 - mitigates the curse of dimensionality, allows for non-local generalisation
 - sparse coding
- Multiple levels of latent variables allow combinatorial sharing of statistical strength

Why should we be bothered with deep learning?

- Effective use of widely available unlabeled data
 - Unsupervised pre-training
 - Semi-supervised learning schemes
- Good performance and efficient solution (expressibility with relatively compact models)
 - better generalisation (lower error on unseen data and lower variance)
 - facilitated optimisation, distinct local minima
 - capacity/complexity control

Why does deep learning seem to work?

- the notion of “*cheap learning*”
 - exponentially fewer parameters than “generic” degrees of freedom (“swindle”)
 - we take advantage of the special nature of problems at hand:
the laws of physics select a particular class of functions that are sufficiently “mathematically simple” to allow “cheap learning” to work
benefitting from *smoothness, symmetry, invariance, locality* (local interactions boosting sparseness)

Henry W. Lin and Max Tegmark, Why does deep and cheap learning work so well?, arXiv:1608.08225


Why does deep learning seem to work?

- the notion of “*cheap learning*”
 - exponentially fewer parameters than “generic” degrees of freedom (“swindle”)
 - we take advantage of the special nature of problems at hand:
the laws of physics select a particular class of functions that are sufficiently “mathematically simple” to allow “cheap learning” to work
benefitting from *smoothness, symmetry, invariance, locality* (local interactions boosting sparseness)
- “*no-flattening*” theorems
 - “flattening polynomials is exponentially expensive, with $2n$ neurons required to multiply n numbers using a single hidden layer, a task that a deep network can perform using only $\sim 4n$ neurons”

Henry W. Lin and Max Tegmark, Why does deep and cheap learning work so well?, arXiv:1608.08225

Why does deep learning seem to work?

- **hierarchical** structure of the physical world
 - hierarchy of the objects and hierarchy of generative processes to untangle
 - decomposition of the generative process into a hierarchy of simpler steps helps reduce the number of parameters (“swindle” paradox)



Henry W. Lin and Max Tegmark, Why does deep and cheap learning work so well?, arXiv:1608.08225

Key challenges ahead

I. Theoretical challenges

- insufficiently tight generalisation bounds (VC dimension)
- difficulty in theoretical handling of complexity of learning in deep architectures ("hard to prove anything")
- is it just another (very efficient) parameterisation of solutions?

II. Visualisation, interpretation, explanation

- explainable deep networks (factors underlying inference outcomes)
- strong initiatives towards visualising and interpreting data representations (particularly in the realm of CNNs)
- how can the process of learning be monitored and controlled?

Key challenges ahead

III. Functionality

- multi-task learning, transfer learning
- multi-modal information processing
- local, incremental learning, self-organisation

IV. Computational challenges

- need for lowering computational costs ("equivalent" networks, performance cost etc.)
- need for better use of data and existing networks (pre-trained)
- dedicated hardware platforms

Summary

- multiple layers facilitate compactness and hierarchy of representations
- deep learning with an originally important role of greedy layer-wise unsupervised pre-training (RBM, AE) – current trend to only rely on supervised learning with ReLU unit activations
- focus on learning multiple levels of representations and optimal discrimination performance
- massive surge of interest fuelled by robust performance of DNNs

Summary and future prospects

- some issues regarding optimal configurations of DNNs , their computational costs and interpretability require further research, especially in light of BIG DATA challenge
- DL concepts bring ML closer to brain inspired computing (models) – promising directions to explore

Summary and future prospects

- some issues regarding optimal configurations of DNNs , their computational costs and interpretability require further research, especially in light of BIG DATA challenge
- DL concepts bring ML closer to brain inspired computing (models) – promising directions to explore

Is it really the direction for machine intelligence in the spirit of general AI?

Recapitulation

- What is the motivation for deep network architectures?
 - expressive power (*expressibility*) and compactness (*efficiency*)
 - hierarchical brain (cortex) organisation
 - multiple levels of abstraction
 - multiple levels of representations suitable for multi-task learning
- *Learning data representations* in deep learning approach vs *hand-engineering features* in traditional pattern recognition
- Learning protocol for DBNs, stacked autoencoders:
 - PHASE I: greedy layer-wise unsupervised pre-training (autoencoders or RBMs)
 - PHASE II: supervised tuning with gradient descent-like optimisation (the last layers or the entire network)

Recapitulation

- Hypotheses about the role of unsupervised pre-training: *regularisation* vs *optimisation* hypotheses
- However, currently there is a trend to avoid pre-training and employ **ReLU units** (less risk for overfitting and local minima)
- What does DL have to offer?
 - learning data representations
 - hierarchy of distributed features (multi-task and transfer learning, non-local generalisation, mitigating the effect and consequences of curse of dimensionality)
 - good performance (large-scale problems) with relatively compact models
 - semi-supervised learning opportunities

Recapitulation

- Why does DL works so well?
 - “cheap learning”
 - “no-flattening” theorems
 - hierarchical structure of the physical work
- Still plenty of challenges ahead!